

Stack Exchange DB

Robert Baranic

```
library(RSQLite)
```

```
list.files()
```

```
## [1] "StackExchange.pdf"      "StackExchangeDB.pdf"    "StackExchangeDB.Rmd"
## [4] "stats.stackexchange.db"
```

```
con = dbConnect(SQLite(), "stats.stackexchange.db")
```

```
sql = function(query) dbGetQuery(con, query)
```

```
dbListTables(con)
```

```
## [1] "BadgeClassMap"      "Badges"                "CloseReasonMap"
## [4] "Comments"          "LinkTypeMap"           "PostHistory"
## [7] "PostHistoryTypeId" "PostLinks"             "PostTypeIdMap"
## [10] "Posts"              "TagPosts"              "Users"
## [13] "VoteTypeMap"        "Votes"
```

1. How many users are there? To do this, we will select the users table and sum the rows since the ID is unique to each user.

```
q = "SELECT COUNT(*)
FROM Users"
sql(q)
```

```
##      COUNT(*)
## 1      321677
```

2. How many users joined since 2020? (Hint: Convert the CreationDate to a year.) To do this, we will convert the CreationDate column to dates and subset the data to >2020. The dates look like this: 2010-07-19T06:55:26.860, so we only need to look at the first 4 characters.

```
q = "SELECT CAST(substr(CreationDate, 1, 4) as decimal) as x
FROM Users
WHERE x >= 2020"
date.df = sql(q)

table(date.df$x)
```

```
##
## 2020 2021 2022 2023
## 34617 32765 28801 4613
```

```
nrow(date.df)
```

```
## [1] 100796
```

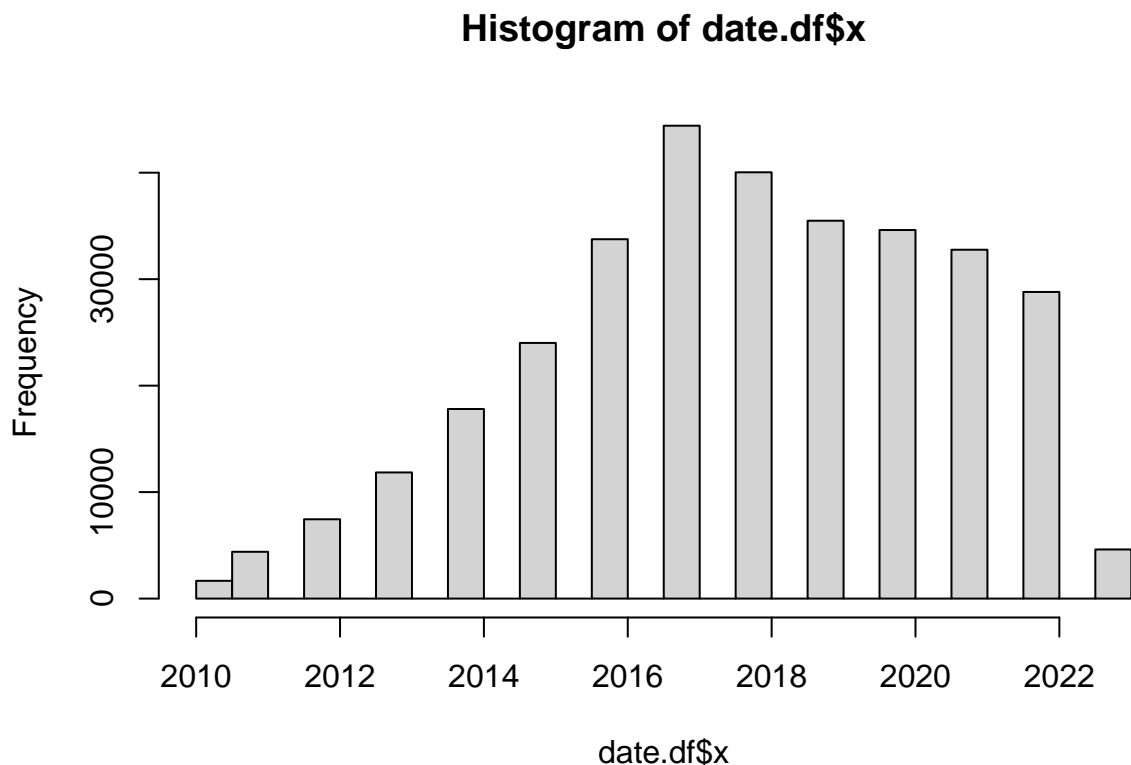
3. How many users joined each year? Describe this with a plot, commenting on any anomalies. We will do this in the same way as before, except we will not subset the data, then we will graph the results. Creation date peaked in 2017 and has been in decline since then. 2023 has very small number of creation date, likely due to it being the current year.

```
q = "SELECT CAST(substr(CreationDate, 1, 4) as decimal) as x
FROM Users"
date.df = sql(q)

table(date.df$x)
```

```
##
## 2010 2011 2012 2013 2014 2015 2016 2017 2018 2019 2020 2021 2022
## 1668 4396 7450 11846 17809 24012 33753 44416 40040 35491 34617 32765 28801
## 2023
## 4613
```

```
hist(date.df$x)
```



- How many different types of posts are there in the Posts table? Get the description of the types from the PostTypeIdMap table. In other words, create a table with the description of each post type and the number of posts of that type, and arrange it from most to least occurrences. Here we will be doing an inner join on the Posts and PostTypeIdMap, joining on their key. Then, we group by PostTypeId and COUNT(*) so we get a sum of the number of those posts. The only other column in PostTypeIdMap is value, so we will join that to the table which contains the description of the post.

```
q = "SELECT PostTypeId, COUNT(*)
FROM Posts
GROUP BY PostTypeId;"

q = "SELECT PostTypeId, Count(*), value
FROM Posts as x
INNER JOIN PostTypeIdMap as y
ON x.PostTypeId = y.id
GROUP BY PostTypeId
ORDER BY COUNT(*) DESC"
sql(q)
```

```
##      PostTypeId Count(*)
## 1              1  204370
## 2              2  197928
## 3              5   1444
## 4              4   1444
## 5              6    23
## 6              3     6
## 7              7     5
##
##                                value
## 1                                Question
## 2                                Answer
## 3                                Tag wiki
## 4                        Tag wiki excerpt
## 5                Moderator nomination
## 6                Orphaned tag wiki
## 7 Wiki placeholder (seems to only be the election description)
```

- How many posted questions are there? From the previous question, we can see that the PostTypeId for “question” is 1, so we can simply subset Posts by this.

```
q = "SELECT COUNT(*)
FROM Posts
WHERE PostTypeId = 1"
sql(q)
```

```
##      COUNT(*)
## 1      204370
```

- What are the top 50 most common tags on questions? For each of the top 50 tags on questions, how many questions are there for each tag. This will be a similar format to #4. We will group Posts by Tags, COUNT(*) and subset where PostId is 1, so it is a question. Then we will order by desc, limit 50, so we get the top 50.

```

q = "SELECT Tags, COUNT(*)
FROM Posts
WHERE PostTypeId = 1
GROUP BY Tags
ORDER BY COUNT(*) DESC
LIMIT 50"
sql(q)

```

##	Tags	COUNT(*)
## 1	<regression>	1019
## 2	<probability>	951
## 3	<hypothesis-testing>	600
## 4	<machine-learning>	435
## 5	<r>	407
## 6	<time-series>	398
## 7	<correlation>	364
## 8	<r><regression>	363
## 9	<neural-networks>	362
## 10	<machine-learning><neural-networks>	324
## 11	<confidence-interval>	290
## 12	<probability><self-study>	284
## 13	<data-visualization>	273
## 14	<regression><logistic>	272
## 15	<clustering>	265
## 16	<pca>	250
## 17	<mathematical-statistics>	241
## 18	<distributions>	234
## 19	<statistical-significance>	233
## 20	<bayesian>	231
## 21	<probability><distributions>	227
## 22	<anova>	205
## 23	<normal-distribution>	195
## 24	<probability><conditional-probability>	193
## 25	<cross-validation>	193
## 26	<multiple-regression>	191
## 27	<regression><multiple-regression>	181
## 28	<sampling>	176
## 29	<time-series><forecasting>	172
## 30	<hypothesis-testing><statistical-significance>	171
## 31	<reinforcement-learning>	169
## 32	<r><time-series>	166
## 33	<r><mixed-model><lme4-nlme>	161
## 34	<t-test>	160
## 35	<mixed-model>	156
## 36	<survival>	152
## 37	<machine-learning><classification>	151
## 38	<clustering><k-means>	141
## 39	<logistic>	139
## 40	<classification>	139
## 41	<time-series><arima>	138
## 42	<self-study>	138
## 43	<meta-analysis>	131
## 44	<chi-squared-test>	128

```
## 45          <regression><correlation>      125
## 46                                <svm>      123
## 47          <standard-deviation>          123
## 48          <markov-process>              123
## 49                                <r><anova>  122
## 50          <machine-learning><cross-validation> 120
```

8. How many answers are there? Like #5, we will use the same form, but use the answer index of 2 instead of 1.

```
q = "SELECT COUNT(*)
FROM Posts
WHERE PostTypeId = 2"
sql(q)
```

```
##  COUNT(*)
## 1    197928
```

9. What's the most recent question (by date-time) in the Posts table? • Find it on the stats.exchange.com Web site and provide the URL. • How would we map a question in the Posts table to the corresponding SO URL? I understood this as the most recent post in terms of creation date, so we will use this column to determine this. We will convert the characters to a date using `datetime()` and order by DESC to get the most recent post. Then, we will also include the `PostId` and the title so we can locate the post to get the URL.

```
q = "SELECT Id, Title, datetime(CreationDate) as Date
FROM Posts
WHERE PostTypeId = 1
ORDER BY Date DESC"

post.df = sql(q)

head(post.df)
```

```
##      Id
## 1 608405
## 2 608403
## 3 608401
## 4 608400
## 5 608398
## 6 608397
##
## 1          Are there any papers or methods that combine mcmc and variational
## 2  An example of  $MA(\infty)$  process with the property of long-term dependence and strict
## 3          What sensitivity analysis technique is suitable for nominal data?
## 4          How sample size affect the chance to reject null hypothesis?
## 5          Is longitudinal data a necessity of mediation paths in structural equation
## 6 Should you impute missing event and time-to-event variables for survival data that has missing and
##      Date
## 1 2023-03-05 05:10:18
## 2 2023-03-05 04:33:18
```

```
## 3 2023-03-05 03:21:49
## 4 2023-03-05 02:51:55
## 5 2023-03-05 01:37:48
## 6 2023-03-05 01:23:04
```

The most recent post was not able to be found, but the second post has the URL as <https://stats.stackexchange.com/questions/608403/an-example-of-ma-infty-process-with-the-property-of-long-term-dependence-and>

We can see a clear pattern as <https://stats.stackexchange.com/questions/%22Id%22/%22title-sepaarated-by-dashes%22> but the title in the URL cuts off at some point. Based on `nchar`, the title is 76 characters with the dashes, which could be the limit. We could map a question from the table to the URL by having <https://stats.stackexchange.com/questions/> followed by the `Id` and the title with `" "` replaced by `"-"` limited to 76 characters, that will get the URL.

```
nchar("an-example-of-ma-infty-process-with-the-property-of-long-term-dependence-and")
```

```
## [1] 76
```

- For the 10 users who posted the most questions • How many questions did they post? • What are the users' names? • When did they join SO? • What is their Reputation? • What country do they have in their profile? To do this, we will have to join the `Posts` table with the `Users` table. We will have to count the number of questions posted by each user and then sort it by `DESC`. We will join them by matching the `Users.Id` to `Posts.OwnerUserId`.

```
q = "SELECT OwnerUserId, DisplayName, u.CreationDate, u.Reputation, u.Location,
      COUNT(DISTINCT p.Id) as Questions
FROM Posts as p, Users as u
WHERE p.OwnerUserId = u.Id
AND PostTypeId = 1
GROUP BY OwnerUserId
ORDER BY Questions DESC
LIMIT 10"

sql(q)
```

##	OwnerUserId	DisplayName	CreationDate	Reputation
## 1	77179	stats_noob	2015-05-14T21:12:31.790	1
## 2	1005	Tim	2010-08-19T15:31:09.537	18497
## 3	9162	user1205901 -	2012-02-13T02:09:08.377	11859
## 4	53690	Richard Hardy	2014-08-08T10:57:13.613	60742
## 5	108150	user321627	2016-03-10T14:45:28.010	2478
## 6	113777	Haitao Du	2016-04-27T20:51:38.203	34665
## 7	28986	Charlie Parker	2013-08-09T19:20:37.540	6286
## 8	40252	An old man in the sea.	2014-02-14T13:18:39.917	5330
## 9	163242	The Pointer	2017-05-30T00:13:41.380	1344
## 10	56211	rnso	2014-09-22T08:35:18.697	9299

##	Location	Questions
## 1		349
## 2		298
## 3		264
## 4	Europe	255

## 5	236
## 6	192
## 7	184
## 8	180
## 9	166
## 10	164

12. For each of the following terms, how many questions contain that term: Regression, ANOVA, Data Mining, Machine Learning, Deep Learning, Neural Network. For this, I interpret it as the title of a question containing any of these terms, not the body. So we will use the Posts table and filter it with LIKE each of the terms, grouping by so using CASE. (source: <https://stackoverflow.com/questions/6101404/sql-group-by-like>)

```
q = "SELECT Title, COUNT(*)
FROM Posts
WHERE (Title LIKE '%Regression%') OR (Title LIKE '%ANOVA%') OR (Title LIKE '%Data
  ↳ Mining%') OR (Title LIKE '%Machine Learning%') OR (Title LIKE '%Deep Learning%') OR
  ↳ (Title LIKE '%Neural Network%')
GROUP BY
  CASE
    WHEN Title LIKE '%Regression%' THEN 'Regression'
    WHEN Title LIKE '%ANOVA%' THEN 'ANOVA'
    WHEN Title LIKE '%Data Mining%' THEN 'Deep Mining'
    WHEN Title LIKE '%Machine Learning%' THEN 'Machine Learning'
    WHEN Title LIKE '%Deep Learning%' THEN 'Deep Learning'
    WHEN Title LIKE '%Neural Network%' THEN 'Neural Network'
    ELSE NULL
  END"

sql(q)
```

##		Title
## 1		How can I adapt ANOVA for binary data?
## 2		Deep learning vs. Decision trees and boosting methods
## 3		Data Mining-- how to tell whether the pattern extracted is meaningful?
## 4		The Two Cultures: statistics vs. machine learning?
## 5		How to choose the number of hidden layers and nodes in a feedforward neural network?
## 6		Can one use multiple regression to predict one principal component (PC) from several other PCs?
##	COUNT(*)	
## 1	2935	
## 2	412	
## 3	116	
## 4	1638	
## 5	2544	
## 6	21886	

I am not sure if SQLite supports the renaming of GROUP BY like the example on SO showed, but this does count the number of posts regardless, just without renaming. The group can be determined by the keyword in the titles.

14. What is the date range for the questions and answers in this database? To do this, we will look at the posts table and convert the creation date to a date and then find the minimum and maximum values for this. Then we will use julianday to do the difference between the two days, but since this is a lot of days, we will divide it 365 to get the number of years difference.

```
q = "SELECT MIN(datetime(CreationDate)), MAX(datetime(CreationDate)),
  ↳ (MAX(julianday(CreationDate)) - MIN(julianday(CreationDate))) / 365
FROM Posts"
```

```
sql(q)
```

```
##      MIN(datetime(CreationDate)) MAX(datetime(CreationDate))
## 1      2009-02-02 14:21:12      2023-03-05 05:10:18
##      (MAX(julianday(CreationDate)) - MIN(julianday(CreationDate))) / 365
## 1                                                                    14.0921
```

15. What question has the most comments associated with it? • how many answers are there for this question? To do this, we will use the posts table and find the post with MAX(CommentCount). From this, we will also show the PostId, Title, and the answer count for this post. We will have to filter for questions again, so we will filter PostTypeId = 1 as known by the earlier command.

```
q = "SELECT Id, Title, MAX(CommentCount), AnswerCount, CreationDate
FROM Posts
WHERE PostTypeId = 1"
```

```
sql(q)
```

```
##      Id
## 1 328630
##
##                                     Title
## 1 Is ridge regression useless in high dimensions ($n \\ll p$)? How can OLS fail to overfit?
##      MAX(CommentCount) AnswerCount      CreationDate
## 1              54              6 2018-02-14T16:31:47.080
```

16. How many comments are there across all posts? • How many posts have a comment? • What is the distribution of comments per question? To do this, we will once again use only the Posts table. I do think we can answer the first two questions in one command, but the third question will need its own command so we can plot it. We will filter the posts such that each post in the table has at least one comment, then we will collapse the table by SUM() and COUNT(*). For the third question, we will filter only question by posttypeid and then we will plot a boxplot of the comment counts.

```
q = "SELECT SUM(CommentCount) as TotalComments, COUNT(*) as NumPostsWithComments
FROM Posts
WHERE CommentCount >= 1"
```


```
sql(q)
```

```
##      TotalComments NumPostsWithComments
## 1          768069          229859
```

```
q = "SELECT CommentCount
FROM Posts
WHERE PostTypeId = 1"
```

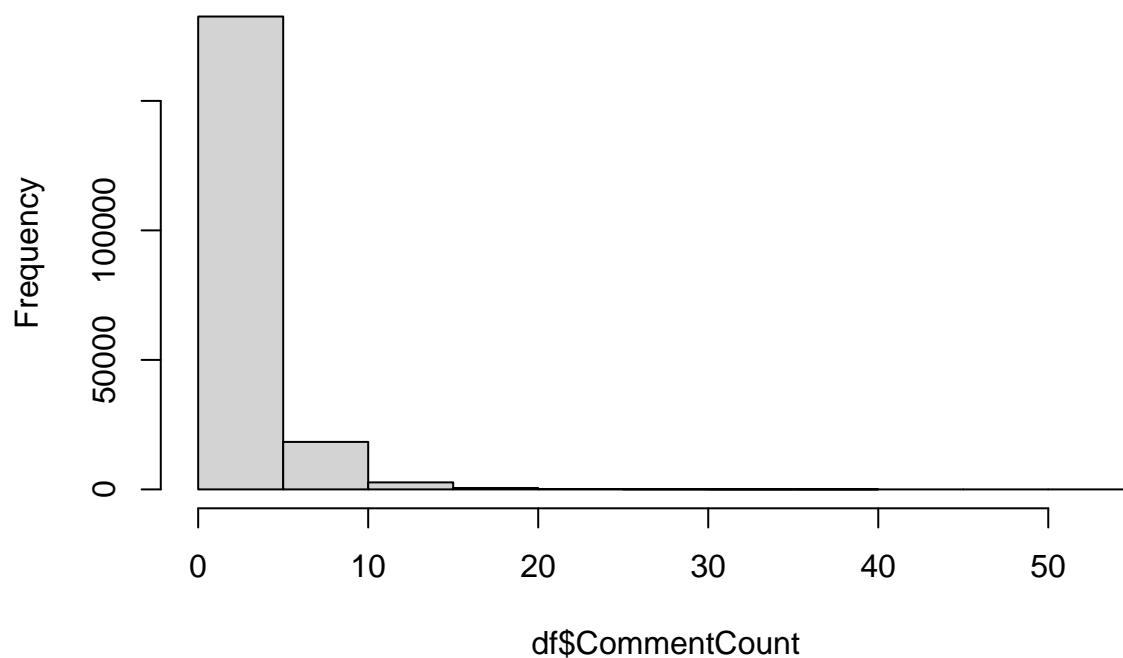
```
df = sql(q)
```

```
boxplot(df$CommentCount, horizontal = T)
```

```
hist(df$CommentCount)
```

Histogram of df\$CommentCount



```
summary(df$CommentCount)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	0.000	0.000	1.000	2.171	3.000	54.000

Required Questions:

21. Compute the table that contains • the question, • the name of the user who posted it, • when that user joined, • their location • the date the question was first posted, • the accepted answer, • when the accepted answer was posted • the name of the user who provided the accepted answer.

To do this, we will have to merge the Posts table with the users table, twice. Once to get the user with the question, and a second time to get the user with the answer so that we can join them both together. I simply went down the list to get each column needed: title of question, display name of question poster, creation date of user, location of user, date of question creation, the body of the accepted answer since these don't have titles, the date of the answer post, and the user who posted it.

```
q = "SELECT q.Title, u1.DisplayName, u1.CreationDate, u1.Location, q.CreationDate,
      a.Body, a.CreationDate, u2.DisplayName
FROM Posts as q, Posts as a, Users as u1, Users as u2
WHERE q.OwnerUserId = u1.Id
AND a.OwnerUserId = u2.Id
AND q.AcceptedAnswerId = a.Id
AND q.PostTypeId = 1"
```

```
head(sql(q))
```

```
##                                     Title
## 1                               Eliciting priors from experts
## 2                               What is normality?
## 3      What are some valuable Statistical Analysis open source projects?
## 4      Assessing the significance of differences in distributions
## 5                               Locating freely available data samples
## 6 Under what conditions should Likert scales be used as ordinal or interval data?
##  DisplayName      CreationDate      Location
## 1 csgillespie 2010-07-19T19:04:52.280 Newcastle, United Kingdom
## 2      A Lion 2010-07-19T19:09:32.157
## 3      grokus 2010-07-19T19:08:29.070 United States
## 4 Jay Stevens 2010-07-19T19:09:16.917 Jacksonville, FL, USA
## 5      EAMann 2010-07-19T19:11:57.393 Tualatin, OR, United States
## 6      A Lion 2010-07-19T19:09:32.157
##      CreationDate
## 1 2010-07-19T19:12:12.510
## 2 2010-07-19T19:12:57.157
## 3 2010-07-19T19:13:28.577
## 4 2010-07-19T19:13:31.617
## 5 2010-07-19T19:15:59.303
## 6 2010-07-19T19:17:47.537
##
## 1
## 2
## 3
## 4
Smirnov test</a>, or the like. The two-sample Kolmogorov-Smirnov test is based on comparing differences
## 5
## 6 <p>Maybe too late but I add my answer anyway...</p>\\\\\\n\\\\\\n<p>It depends on what you intend to do
##      CreationDate      DisplayName
## 1 2010-07-19T19:19:46.160 Harlan
## 2 2010-07-19T19:43:20.423 John L. Taylor
## 3 2010-07-19T19:14:43.050 Jay Stevens
```

```
## 4 2010-07-19T21:36:12.850 John L. Taylor
## 5 2010-07-19T19:24:18.580 Stephen Turner
## 6 2010-08-19T10:00:00.370 chl
```

22. Determine the users that have only posted questions and never answered a question? (Compute the table containing the number of questions, number of answers and the user's login name for this group.) How many are there? To do this, I will have to join the Posts and Users table once again. We will join on OwnerUserId = Id, then we will have to find the frequency of posts and answers. To do this, I found the most efficient way is to use SUM(CASE...) where we will have a case when PostTypeID = 1 and 2, so that the sum will be the frequency (source: <https://stackoverflow.com/questions/24767130/get-frequency-of-a-column-in-sql-server>). Then we can order by the number of answers to see those who have not answered. Then, to find those who have posted at least 1 question, I nested that query and subsetting it from there. I found it easiest to just table the number of answers to find the number of those who have not answered.

```
q = "SELECT Questions, Answers
FROM (

SELECT u.Id, u.DisplayName, SUM(CASE WHEN PostTypeId = 1 then 1 else 0 end) as Questions,
  SUM(CASE WHEN PostTypeId = 2 then 1 else 0 end) as Answers
FROM Posts as p, Users as u
WHERE p.OwnerUserId = u.Id
GROUP BY u.Id
ORDER BY Answers) as innerquery
WHERE Questions > 0"

df = sql(q)
head(df)
```

```
##   Questions Answers
## 1         2        0
## 2         2        0
## 3         1        0
## 4        13        0
## 5         2        0
## 6         1        0
```

```
head(table(df$Answers))
```

```
##
##      0      1      2      3      4      5
## 76410 6470 1864  888  557  372
```

23. Compute the table with information for the 75 users with the most accepted answers. This table should include • the user's display name, • creation date, • location, • the number of badges they have won, – the names of the badges (as a single string) • the dates of the earliest and most recent accepted answer (as two fields) – the (unique) tags for all the questions for which they had the accepted answer (as a single string) Due to limitations of my machine, I could not get a decent runtime while using GROUP_CONCAT(b.Name) to put all badge names into one string, but this did work when I used LIMIT = 1, just to experiment, even then, it made R studio unbearable.

To do this question, it required a total of 4 tables, all left joins since this made it much simpler and quicker to run. I used Id from the Users table as the primary key that linked all the tables together and I grouped them on this Id. Users table contained all the information on the user that we needed. However, first we needed to calculate the top 75 accepted answers users, so we linked the posts table twice, one to the question and one to the accepted answer that linked back to the users table. From that, we could simply COUNT the distinct Id from the answers table so we could order by DESC. Then we included the badges table where we simply linked the Id's together and could COUNT Id from the badges table. As prefaced above, group_concat for the badge names would not run within a reasonable amount of time, but I did manage to group_concat for the tags since they were a shorter string. Min and Max datetime() were added too.

```
q = "SELECT u.Id, u.DisplayName, u.CreationDate, COUNT(DISTINCT a.Id) as AcceptedAnswers,
  ↳ COUNT(DISTINCT b.Id) as NumBadges, MIN(datetime(a.CreationDate)) as FirstAnswer,
  ↳ MAX(datetime(a.CreationDate)) as LatestAnswer, GROUP_CONCAT(DISTINCT q.Tags) as
  ↳ QuestionTags
FROM Users as u
LEFT JOIN Posts as q
ON q.OwnerUserID = u.Id
LEFT JOIN Posts as a
ON a.Id = q.AcceptedAnswerId
LEFT JOIN Badges as b
ON b.UserId = u.Id
GROUP BY u.Id
ORDER BY AcceptedAnswers DESC
LIMIT 75"
```

```
df = sql(q)
head(df)
```

##	Id	DisplayName	CreationDate	AcceptedAnswers
## 1	9162	user1205901 -	2012-02-13T02:09:08.377	179
## 2	108150	user321627	2016-03-10T14:45:28.010	173
## 3	163242	The Pointer	2017-05-30T00:13:41.380	124
## 4	253	Tal Galili	2010-07-21T07:53:50.990	115
## 5	56211	rnso	2014-09-22T08:35:18.697	112
## 6	1005	Tim	2010-08-19T15:31:09.537	110

##	NumBadges	FirstAnswer	LatestAnswer
## 1	275	2012-05-16 05:36:25	2023-02-19 16:13:21
## 2	87	2016-03-10 15:11:46	2023-02-13 03:29:21
## 3	66	2017-06-06 10:41:22	2023-01-20 18:29:46
## 4	369	2010-08-09 18:00:24	2021-11-15 15:50:42
## 5	186	2014-10-01 09:50:21	2022-02-21 12:02:21
## 6	353	2010-09-30 16:32:06	2014-12-08 23:54:00

## 1	<multilevel-analysis><dyadic-data>,<forecasting><error><mse><mae>,<forecasting>,<confidence
## 2	<mathematical-statistics><asymptotics><umvue>,<confidence-interval>,<mathematical-statistics><suff
## 3	
## 4	
## 5	
## 6	

24. How many questions received no answers (accepted or unaccepted)? How many questions had no accepted answer? To do this in a simpler way, I broke it into two sql commands, one for the first question and one for the second. I used the same format for both, but on the first command, I linked

the two Posts tables with the parentId of the answer with the question Id. Then, since this was a left join, I could sum over the NULL values if the question did not have an answer. Same process for the accepted answer except that I linked the accepted answer id from the question to the id of the answer post.

```
q = "SELECT SUM(CASE WHEN a.id IS NULL THEN 1 ELSE 0 END) as NoAnswer
FROM Posts as q
LEFT JOIN Posts as a
ON a.ParentId = q.Id
WHERE q.PostTypeId = 1"

sql(q)
```

```
##      NoAnswer
## 1      66970
```

No Accepted Answers:

```
q = "SELECT SUM(CASE WHEN a.id IS NULL THEN 1 ELSE 0 END) as NoAcceptedAnswer
FROM Posts as q
LEFT JOIN Posts as a
ON q.AcceptedAnswerId = a.Id
WHERE q.PostTypeId = 1"

sql(q)
```

```
##      NoAcceptedAnswer
## 1      136366
```


25. What is the distribution of answers per posted question? For this, I tabled the answercount from the Posts table. From this I compared it to when I left joined Posts and Posts together on Id and ParentId. When I summed the number of answers, I could only get it to sum when it had one or more, ie it would not return 0 if a post did not have any answers. When comparing to the table of the first command, the numbers are the same except the first table has a count for 0 answers and the second does not. The distribution is very heavily weighted towards 0-1-2 answers and very few greater than that.

```
q = "SELECT q.Id, q.AnswerCount
FROM Posts as q
WHERE q.PostTypeId = 1"

df = sql(q)
table(df$AnswerCount)
```

```
##
##      0      1      2      3      4      5      6      7      8      9     10     11     12
## 66970 98602 27191 7246 2408  905  401  210  136  82   62   35   25
##    13    14    15    16    17    18    19    20    21    22    23    24    25
##    19    14    14     8     5     4     5     1     5     2     1     1     2
##    26    27    28    31    32    34    37    38    46    47    78    80   153
##     2     1     2     1     2     1     1     1     1     1     1     1     1
```

```
boxplot(df$AnswerCount, horizontal = T)
```



```
summary(df$AnswerCount)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
##  0.0000   0.0000   1.0000   0.9685   1.0000  153.0000
```

```
q = "SELECT q.Id, SUM(CASE WHEN a.Id IS NOT NULL THEN 1 ELSE 0 END)
FROM Posts as q
LEFT JOIN Posts as a
ON a.ParentId = q.Id
WHERE a.PostTypeId = 2
AND q.PostTypeId = 1
GROUP BY q.Id"
```

```
df = sql(q)
```

```
table(df$`SUM(CASE WHEN a.Id IS NOT NULL THEN 1 ELSE 0 END)`)
```

```
##
##      1      2      3      4      5      6      7      8      9     10     11     12     13
## 98602 27191 7246 2408  905  401  210  136   82   62   35   25   19
##     14     15     16     17     18     19     20     21     22     23     24     25     26
##     14     14      8      5      4      5      1      5      2      1      1      2      2
##     27     28     31     32     34     37     38     46     47     78     80    153
##      1      2      1      2      1      1      1      1      1      1      1      1
```

26. What is the length of time for a question to receive an answer? to obtaining an accepted answer? To do this, we will join the posts table with itself on `Id = ParentId` for all questions and `Id = AcceptedAnswerId` for the accepted answers. When this is done, we can see that there are some outliers where the answer is received before the question, but in most cases, the answer comes after the questions.

```

q = "SELECT q.Id, q.Title, q.CreationDate as QuestionDate, a.CreationDate as AnswerDate,
      ↪ julianday(a.CreationDate) - julianday(q.CreationDate) as AnswerTime
FROM Posts as q, Posts as a
WHERE q.Id = a.ParentId
AND q.PostTypeId = 1
AND a.PostTypeId = 2
ORDER BY AnswerTime"

df = sql(q)
head(df)

```

```

##      Id                                     Title
## 1 112451      Maximum Likelihood Estimation (MLE) in layman terms
## 2 112451      Maximum Likelihood Estimation (MLE) in layman terms
## 3  16921      How to understand degrees of freedom?
## 4  16921      How to understand degrees of freedom?
## 5 116804 Coefficient changes sign when adding a variable in logistic regression
## 6 239431      Book recommendations for probability
##      QuestionDate      AnswerDate AnswerTime
## 1 2014-08-19T12:46:29.950 2012-12-03T23:52:18.377 -623.5376
## 2 2014-08-19T12:46:29.950 2012-12-04T13:09:31.343 -622.9840
## 3 2011-10-12T20:16:52.280 2010-07-28T12:48:14.233 -441.3116
## 4 2011-10-12T20:16:52.280 2010-07-28T12:49:31.780 -441.3107
## 5 2014-09-25T21:49:50.467 2013-12-16T21:13:01.827 -283.0256
## 6 2016-10-10T15:24:14.047 2016-01-11T15:19:13.797 -273.0035


```

```

boxplot(df$AnswerTime, horizontal = T, xlab = "Days", main = "Distribution of Answer
      ↪ Times")

```

Distribution of Answer Times



```
summary(df$AnswerTime)
```

```
##      Min.   1st Qu.   Median     Mean  3rd Qu.    Max.
## -623.538    0.047    0.247   145.172    3.857  4578.530
```


```
q = "SELECT q.Id, q.Title, q.CreationDate as QuestionDate, a.CreationDate as AnswerDate,
  ↪ julianday(a.CreationDate) - julianday(q.CreationDate) as AnswerTime
FROM Posts as q, Posts as a
WHERE q.AcceptedAnswerId = a.Id
AND q.PostTypeId = 1
AND a.PostTypeId = 2
ORDER BY AnswerTime"
```

```
df = sql(q)
head(df)
```

```
##      Id                                     Title
## 1 129091                               Appropriate statistical analysis to use
## 2 108295                               Algorithm for rating books: Relative perception
## 3 34165                                Expected value of the log-determinant of a Wishart matrix
## 4 51070    How can I convert a q-value distribution to a p-value distribution?
## 5 63391    Free internet or downloadable resources for sample size calculations
## 6 64538    How do I find values not given in (interpolate in) statistical tables?
##      QuestionDate      AnswerDate AnswerTime
## 1 2014-12-14T20:27:18.950 2014-12-14T18:41:59.110 -0.07314630
## 2 2014-07-17T11:36:45.187 2014-07-17T11:09:39.077 -0.01882072
## 3 2012-08-12T03:36:55.910 2012-08-12T03:36:55.910 0.00000000
## 4 2013-02-28T20:05:01.313 2013-02-28T20:05:01.313 0.00000000
## 5 2013-07-05T04:10:52.657 2013-07-05T04:10:52.657 0.00000000
## 6 2013-07-16T23:55:34.533 2013-07-16T23:55:34.533 0.00000000
```

```
boxplot(df$AnswerTime, horizontal = T, xlab = "Days", main = "Distribution of Accepted
  ↪ Answer Times")
```


Distribution of Accepted Answer Times



```
summary(df$AnswerTime)
```

```
##      Min.   1st Qu.   Median     Mean  3rd Qu.     Max.
## -0.073    0.036    0.141   25.306   0.919 3877.244
```

27. How many answers are typically received before the accepted answer? To do this, I first did this individually for each question, then found the average in another command after I knew this would work. I joined Posts three times together, once for the question, second time for all the answers, and third time for the accepted answer. From this, I used SUM(CASE) once again such that it summed all the times when the creation date of the answer not accepted was less than the creation date of the accepted answer. From there, I averaged all of the values in another sql command by using the first command as a nested query. I did mean() in R to validate the results.

```
q = "SELECT q.Id, q.Title, SUM(CASE WHEN julianday(a.CreationDate) <
  ↳ julianday(ac.CreationDate) THEN 1 ELSE 0 END) as NumAnswersBeforeAccepted
FROM Posts as q, Posts as a, Posts as ac
WHERE q.Id = a.ParentId
AND q.AcceptedAnswerId = ac.Id
AND q.PostTypeId = 1
AND a.PostTypeId = 2
AND ac.PostTypeId = 2
GROUP BY q.Id"
df = sql(q)
head(df)
```

```
##      Id
## 1     1
```

```
## 2 2
## 3 3
## 4 4
## 5 7
## 6 10
##
##                                     Title
## 1                                     Eliciting priors from experts
## 2                                     What is normality?
## 3          What are some valuable Statistical Analysis open source projects?
## 4          Assessing the significance of differences in distributions
## 5                                     Locating freely available data samples
## 6 Under what conditions should Likert scales be used as ordinal or interval data?
##   NumAnswersBeforeAccepted
## 1                         0
## 2                         1
## 3                         0
## 4                         1
## 5                         1
## 6                         2
```

```
mean(df$NumAnswersBeforeAccepted)
```

```
## [1] 0.1836068
```

```
q = "SELECT AVG(NumAnswersBeforeAccepted)
FROM
(SELECT q.Id, q.Title, SUM(CASE WHEN julianday(a.CreationDate) <
  ↳ julianday(ac.CreationDate) THEN 1 ELSE 0 END) as NumAnswersBeforeAccepted
FROM Posts as q, Posts as a, Posts as ac
WHERE q.Id = a.ParentId
AND q.AcceptedAnswerId = ac.Id
AND q.PostTypeId = 1
AND a.PostTypeId = 2
AND ac.PostTypeId = 2
GROUP BY q.Id) as innerquery"

sql(q)
```

```
##   AVG(NumAnswersBeforeAccepted)
## 1                         0.1836068
```